

## Question 1: Load

Programmatically download and load into your favorite analytical tool the transactions data. This data, which is in line-delimited JSON format, can be found [here](#)

**Please describe the structure of the data. Number of records and fields in each record?**


There are 786363 transactions in the JSON data. It is list of dictionaries. There are a total of 786363 transaction records. For each record, there are 29 fields.

**Please provide some additional basic summary statistics for each field. Be sure to include a count of null, minimum, maximum, and unique values where appropriate.**

1. There are 4562 missing data in acqCountry, 724 in merchantCountryCode, 4054 in posEntryMode, 409 in posConditionCode, 698 in transactionType. There are no data for these 6 categories: echoBuffer, merchantCity, merchantState, merchantZip, posOnPremises, recurringAuthInd.
2. The range for creditLimit is from 250 to 50000. The range for availableMoney is from -1005.63 to 50000. The range for transactionAmount is from 0 to 2011.54. The range for currentBalance is from 0 to 47498.81.
3. There are a total of 5000 unique customer ID in this dataset. For these customers, there are a total of 4 countries recorded (not including NA). There are 19 merchantCategoryCode types, and 3 different transaction types (not including NA).

## Question 2: Plot

**Plot a histogram of the processed amounts of each transaction, the transactionAmount column.**



(histogram for amount of transaction)

**Report any structure you find and any hypotheses you have about that structure.**

Exponential decay. Number of transactions decrease exponentially as the transaction amount increases

## Question 3: Data Wrangling - Duplicate Transactions

You will notice a number of what look like duplicated transactions in the data set. One type of duplicated transaction is a reversed transaction, where a purchase is followed by a reversal. Another example is a multi-swipe, where a vendor accidentally charges a customer's card multiple times within a short time span.

**Can you programmatically identify reversed and multi-swipe transactions?**

Reversed transactions are recorded in `reverse_df`. They are identified as they have `transactionType` 'REVERSAL'.

Multi-swipe transactions are recorded in `Multi_df_f2`. They are identified as multiple transactions of the same amount, same customer, same merchant, within 2 minutes.

**What total number of transactions and total dollar amount do you estimate for the reversed transactions? For the multi-swipe transactions? (please consider the first transaction to be "normal" and exclude it from the number of transaction and dollar amount counts)**

There are a total of 20303 reversed transactions. The total amount for the reversed transactions is 2821792.5.

There are a total of 666 multi-swipe transactions. The total amount for the multi-swipe transactions is 40266.25.

**Did you find anything interesting about either kind of transaction?**


There are a total of 19 types of `merchant_Category`. Online retail are the `merchant_Category` that most subject to mistakes.

For reversed transactions, there are only 13 types involved. The most common type is online retail (28%) and fastfood (16%). There are no following types of reversed transactions: 'cable/phone', 'food\_delivery', 'fuel', 'gym', 'mobileapps', 'online\_subscriptions'. Also, there are higher proportions of transactions in the range of 0-200 compared to that of multi-swipe.



(Reversed Transactions Pie Chart)

For multi-swipe transactions, all 19 types involved. The most common type is online mobileApps (24%) and online subscription (17%). The two most common type are both online payments.



(Multi-swipe Transactions Pie Chart)

## Question 4: Model


Fraud is a problem for any bank. Fraud can take many forms, whether it is someone stealing a single credit card, to large batches of stolen credit card numbers being used on the web, or even a mass compromise of credit card numbers stolen from a merchant via tools like credit card skimming devices.

Each of the transactions in the dataset has a field called isFraud. Please build a predictive model to determine whether a given transaction will be fraudulent or not. Use as much of the data as you like (or all of it).

**Provide an estimate of performance using an appropriate sample, and show your work.**

Because a very small proportion of data is positive, I used AUCROC score to evaluate the performance. The data use split at 0.8/0.2 for training/test. Comparing the performance of SVC, Random forest, logistic regression, adaboost and xgboost, it is found that xgboost gives the best AUC of 0.99. The accuracy of xgboost is also the highest among algorithms.

Please explain your methodology (modeling algorithm/method used and why, what features/data you found useful, what questions you have, and what you would do next with more time)




(AUC for xgboost)

1. Since this is a binary prediction for fraud, I used binary Classification models including SVC, Random forest logistic regression, adaboost and xgboost, it is found that xgboost gives both the best AUC score of 0.99 and the best accuracy of 0.95.
2. Using a parameter analysis on feature importance, the creditLimit data was excluded due to low feature importance. Some of the top important features are: merchantCountryCode\_US, cardPresent\_True, merchantCategoryCode\_online\_retail. The following variable is included: 'availableMoney', 'transactionDateTime' (transformed to months and hours), 'merchantCountryCode', 'transactionAmount', 'merchantCategoryCode', 'transactionType', 'currentBalance', 'cardPresent', 'expirationDateKeyInMatch', 'isFraud'

	<b>features</b>	<b>coef</b>	<b>abs_importance</b>
<b>25</b>	merchantCountryCode_US	-1.442520	1.442520
<b>27</b>	transactionType_PURCHASE	-1.353383	1.353383
<b>28</b>	cardPresent_True	-0.718365	0.718365
<b>17</b>	merchantCategoryCode_online_retail	-0.264183	0.264183
<b>7</b>	merchantCategoryCode_fastfood	-0.222783	0.222783
<b>8</b>	merchantCategoryCode_food	-0.167477	0.167477
<b>6</b>	merchantCategoryCode_entertainment	-0.161233	0.161233
<b>26</b>	transactionType_ADDRESS_VERIFICATION	-0.098699	0.098699
<b>10</b>	merchantCategoryCode_fuel	-0.097059	0.097059
<b>16</b>	merchantCategoryCode_online_gifts	-0.091584	0.091584
<b>15</b>	merchantCategoryCode_mobileapps	-0.077095	0.077095
<b>20</b>	merchantCategoryCode_rideshare	-0.065211	0.065211
<b>18</b>	merchantCategoryCode_online_subscriptions	-0.056763	0.056763
<b>4</b>	merchantCategoryCode_auto	-0.048270	0.048270
<b>14</b>	merchantCategoryCode_hotels	-0.045530	0.045530
<b>13</b>	merchantCategoryCode_health	-0.039057	0.039057
<b>21</b>	merchantCategoryCode_subscriptions	-0.033375	0.033375
<b>19</b>	merchantCategoryCode_personal care	-0.032459	0.032459
<b>9</b>	merchantCategoryCode_food_delivery	-0.021027	0.021027
<b>11</b>	merchantCategoryCode_furniture	-0.011372	0.011372

(list of feature importance)

1. Because there are much more non-fraud cases compared to fraud cases, the more abundant categories seem to have a greater negative impact. A more balanced dataset, i.e a greater proportion of fraud cases included, may change the feature importance and increase prediction rate.



(bar graph of total number of NotFraud and Fraud cases)

1. If given more time, I wish to perform clustering analysis to cluster customers into groups by features, and see if each groups have different fraud rate and prediction accuracy.

## Code

```
In [1]: import json
import pandas as pd
import numpy as np
import datetime as dt
```

```
In [2]: transactions = []
for line in open('transactions.json', 'r'):
    transactions.append(json.loads(line))
```

```
In [3]: transactions[0:3]
```

```
Out[3]: [{'accountNumber': '737265056',
'customerId': '737265056',
'creditLimit': 5000.0,
'availableMoney': 5000.0,
'transactionDateTime': '2016-08-13T14:27:32',
'transactionAmount': 98.55,
'merchantName': 'Uber',
'acqCountry': 'US',
'merchantCountryCode': 'US',
'posEntryMode': '02',
'posConditionCode': '01',
'merchantCategoryCode': 'rideshare',
'currentExpDate': '06/2023',
'accountOpenDate': '2015-03-14',
'dateOfLastAddressChange': '2015-03-14',
'cardCVV': '414',
'enteredCVV': '414',
'cardLast4Digits': '1803',
'transactionType': 'PURCHASE',
'echoBuffer': ''},
```

```
'currentBalance': 0.0,
'merchantCity': '',
'merchantState': '',
'merchantZip': '',
'cardPresent': False,
'posOnPremises': '',
'recurringAuthInd': '',
'expirationDateKeyInMatch': False,
'isFraud': False},
{ 'accountNumber': '737265056',
  'customerId': '737265056',
  'creditLimit': 5000.0,
  'availableMoney': 5000.0,
  'transactionDateTime': '2016-10-11T05:05:54',
  'transactionAmount': 74.51,
  'merchantName': 'AMC #191138',
  'acqCountry': 'US',
  'merchantCountryCode': 'US',
  'posEntryMode': '09',
  'posConditionCode': '01',
  'merchantCategoryCode': 'entertainment',
  'cardPresent': True,
  'currentExpDate': '02/2024',
  'accountOpenDate': '2015-03-14',
  'dateOfLastAddressChange': '2015-03-14',
  'cardCVV': '486',
  'enteredCVV': '486',
  'cardLast4Digits': '767',
  'transactionType': 'PURCHASE',
  'echoBuffer': '',
  'currentBalance': 0.0,
  'merchantCity': '',
  'merchantState': '',
  'merchantZip': '',
  'posOnPremises': '',
  'recurringAuthInd': '',
  'expirationDateKeyInMatch': False,
  'isFraud': False},
{ 'accountNumber': '737265056',
  'customerId': '737265056',
  'creditLimit': 5000.0,
  'availableMoney': 5000.0,
  'transactionDateTime': '2016-11-08T09:18:39',
  'transactionAmount': 7.47,
  'merchantName': 'Play Store',
  'acqCountry': 'US',
  'merchantCountryCode': 'US',
  'posEntryMode': '09',
  'posConditionCode': '01',
  'merchantCategoryCode': 'mobileapps',
  'currentExpDate': '08/2025',
  'accountOpenDate': '2015-03-14',
  'dateOfLastAddressChange': '2015-03-14',
  'cardCVV': '486',
  'enteredCVV': '486',
  'cardLast4Digits': '767',
  'transactionType': 'PURCHASE',
  'echoBuffer': '',
  'currentBalance': 0.0,
  'merchantCity': '',
  'merchantState': '',
  'merchantZip': '',
  'cardPresent': False,
  'posOnPremises': '',
  'recurringAuthInd': ''}
```

```
'expirationDateKeyInMatch': False,
'isFraud': False}]
```

```
In [4]: len(transactions)
```

```
Out[4]: 786363
```

```
In [5]: print(len(transactions[1].keys()))
```

```
29
```

```
In [6]: df = pd.DataFrame(transactions)
```

```
In [7]: df.head(20)
```

```
Out[7]:
```

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transactionAmount
0	737265056	737265056	5000.0	5000.00	2016-08-13T14:27:32	9
1	737265056	737265056	5000.0	5000.00	2016-10-11T05:05:54	7
2	737265056	737265056	5000.0	5000.00	2016-11-08T09:18:39	
3	737265056	737265056	5000.0	5000.00	2016-12-10T02:14:50	
4	830329091	830329091	5000.0	5000.00	2016-03-24T21:04:46	;
5	830329091	830329091	5000.0	5000.00	2016-04-19T16:24:27	3
6	830329091	830329091	5000.0	5000.00	2016-05-21T14:50:35	5
7	830329091	830329091	5000.0	5000.00	2016-06-03T00:31:21	
8	830329091	830329091	5000.0	4990.63	2016-06-10T01:21:46	52
9	830329091	830329091	5000.0	5000.00	2016-07-11T10:47:16	16
10	830329091	830329091	5000.0	5000.00	2016-09-07T20:22:47	16
11	830329091	830329091	5000.0	5000.00	2016-12-07T16:34:04	4
12	830329091	830329091	5000.0	4959.25	2016-12-14T10:00:35	4
13	830329091	830329091	5000.0	4918.50	2016-12-20T18:38:23	4
14	830329091	830329091	5000.0	4877.75	2016-12-28T06:43:01	4
15	574788567	574788567	2500.0	2500.00	2016-01-02T11:19:46	3
16	574788567	574788567	2500.0	2469.92	2016-01-16T01:01:27	4
17	574788567	574788567	2500.0	2428.67	2016-01-26T14:04:22	



	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transactionAmount
18	574788567	574788567	2500.0	2428.67	2016-01-29T07:17:39	12
19	574788567	574788567	2500.0	2304.46	2016-01-29T07:33:15	19

20 rows x 29 columns

In [8]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 786363 entries, 0 to 786362
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   accountNumber                        786363 non-null object
1   customerId                          786363 non-null object
2   creditLimit                         786363 non-null float64
3   availableMoney                      786363 non-null float64
4   transactionDateTime                 786363 non-null object
5   transactionAmount                  786363 non-null float64
6   merchantName                      786363 non-null object
7   acqCountry                        786363 non-null object
8   merchantCountryCode                786363 non-null object
9   posEntryMode                      786363 non-null object
10  posConditionCode                   786363 non-null object
11  merchantCategoryCode               786363 non-null object
12  currentExpDate                     786363 non-null object
13  accountOpenDate                    786363 non-null object
14  dateOfLastAddressChange            786363 non-null object
15  cardCVV                           786363 non-null object
16  enteredCVV                         786363 non-null object
17  cardLast4Digits                    786363 non-null object
18  transactionType                    786363 non-null object
19  echoBuffer                         786363 non-null object
20  currentBalance                     786363 non-null float64
21  merchantCity                       786363 non-null object
22  merchantState                      786363 non-null object
23  merchantZip                        786363 non-null object
24  cardPresent                        786363 non-null bool
25  posOnPremises                      786363 non-null object
26  recurringAuthInd                   786363 non-null object
27  expirationDateKeyInMatch           786363 non-null bool
28  isFraud                           786363 non-null bool
dtypes: bool(3), float64(4), object(22)
memory usage: 158.2+ MB
```

In [9]:

```
df.describe()
```

Out[9]:

	creditLimit	availableMoney	transactionAmount	currentBalance
<b>count</b>	786363.000000	786363.000000	786363.000000	786363.000000
<b>mean</b>	10759.464459	6250.725369	136.985791	4508.739089
<b>std</b>	11636.174890	8880.783989	147.725569	6457.442068
<b>min</b>	250.000000	-1005.630000	0.000000	0.000000
<b>25%</b>	5000.000000	1077.420000	33.650000	689.910000

	creditLimit	availableMoney	transactionAmount	currentBalance
<b>50%</b>	7500.000000	3184.860000	87.900000	2451.760000
<b>75%</b>	15000.000000	7500.000000	191.480000	5291.095000
<b>max</b>	50000.000000	50000.000000	2011.540000	47498.810000

```
In [10]: #check duplicates
len(df)-len(df.drop_duplicates())
```

```
Out[10]: 0
```

```
In [11]: df.columns
```

```
Out[11]: Index(['accountNumber', 'customerId', 'creditLimit', 'availableMoney',
               'transactionDateTime', 'transactionAmount', 'merchantName',
               'acqCountry', 'merchantCountryCode', 'posEntryMode', 'posConditionCode',
               'merchantCategoryCode', 'currentExpDate', 'accountOpenDate',
               'dateOfLastAddressChange', 'cardCVV', 'enteredCVV', 'cardLast4Digits',
               'transactionType', 'echoBuffer', 'currentBalance', 'merchantCity',
               'merchantState', 'merchantZip', 'cardPresent', 'posOnPremises',
               'recurringAuthInd', 'expirationDateKeyInMatch', 'isFraud'],
              dtype='object')
```

```
In [123... #check for unique values
df.nunique()
```

```
Out[123... accountNumber      5000
customerId      5000
creditLimit      10
availableMoney  521916
transactionDateTime  776637
transactionAmount  66038
merchantName     2490
acqCountry        5
merchantCountryCode  5
posEntryMode      6
posConditionCode   4
merchantCategoryCode  19
currentExpDate    165
accountOpenDate   1820
dateOfLastAddressChange  2184
cardCVV           899
enteredCVV        976
cardLast4Digits   5246
transactionType    4
echoBuffer         1
currentBalance   487318
merchantCity       1
merchantState      1
merchantZip        1
cardPresent        2
posOnPremises      1
recurringAuthInd   1
expirationDateKeyInMatch  2
isFraud            2
dtype: int64
```

```
In [13]: df.posConditionCode.unique()
```

```
Out[13]: array(['01', '08', '99', ''], dtype=object)
```

```
In [14]: df.posEntryMode.unique()
```

```
Out[14]: array(['02', '09', '05', '80', '90', ''], dtype=object)
```

```
In [15]: df.transactionType.unique()
```

```
Out[15]: array(['PURCHASE', 'ADDRESS_VERIFICATION', 'REVERSAL', ''], dtype=object)
```

```
In [16]: df2=df.copy()
df2.acqCountry[8]
df2=df2.replace(' ', np.nan)
df2.acqCountry[8]
```

```
Out[16]: nan
```

```
In [17]: #check for null
df2.isna().sum()
```

```
Out[17]: accountNumber      0
customerId      0
creditLimit      0
availableMoney    0
transactionDateTime  0
transactionAmount  0
merchantName      0
acqCountry      4562
merchantCountryCode  724
posEntryMode      4054
posConditionCode   409
merchantCategoryCode  0
currentExpDate      0
accountOpenDate      0
dateOfLastAddressChange  0
cardCVV            0
enteredCVV          0
cardLast4Digits      0
transactionType      698
echoBuffer          786363
currentBalance        0
merchantCity          786363
merchantState          786363
merchantZip           786363
cardPresent           0
posOnPremises          786363
recurringAuthInd       786363
expirationDateKeyInMatch  0
isFraud               0
dtype: int64
```

```
In [18]: df2.columns
#df2=df2.loc[:,['accountNumber', 'customerId', 'creditLimit', 'availableMoney',
#              'transactionDateTime', 'transactionAmount', 'merchantName',
```

```

#         'acqCountry', 'merchantCountryCode',
#         'merchantCategoryCode', 'currentExpDate', 'accountOpenDate',
#         'dateOfLastAddressChange',
#         'transactionType', 'currentBalance', 'cardPresent', 'expirationDateKeyIn

df2=df2.loc[:,['accountNumber', 'customerId', 'creditLimit', 'availableMoney',
               'transactionDateTime', 'transactionAmount',
               'acqCountry', 'merchantCountryCode',
               'merchantCategoryCode', 'currentExpDate', 'accountOpenDate',
               'dateOfLastAddressChange', 'cardCVV', 'enteredCVV',
               'transactionType', 'currentBalance', 'cardPresent', 'expirationDateKeyIn
df2.head(50)

```

Out[18]:

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transactionAm
0	737265056	737265056	5000.0	5000.00	2016-08-13T14:27:32	9
1	737265056	737265056	5000.0	5000.00	2016-10-11T05:05:54	.
2	737265056	737265056	5000.0	5000.00	2016-11-08T09:18:39	
3	737265056	737265056	5000.0	5000.00	2016-12-10T02:14:50	
4	830329091	830329091	5000.0	5000.00	2016-03-24T21:04:46	
5	830329091	830329091	5000.0	5000.00	2016-04-19T16:24:27	3
6	830329091	830329091	5000.0	5000.00	2016-05-21T14:50:35	5
7	830329091	830329091	5000.0	5000.00	2016-06-03T00:31:21	
8	830329091	830329091	5000.0	4990.63	2016-06-10T01:21:46	52
9	830329091	830329091	5000.0	5000.00	2016-07-11T10:47:16	16
10	830329091	830329091	5000.0	5000.00	2016-09-07T20:22:47	16
11	830329091	830329091	5000.0	5000.00	2016-12-07T16:34:04	4
12	830329091	830329091	5000.0	4959.25	2016-12-14T10:00:35	4
13	830329091	830329091	5000.0	4918.50	2016-12-20T18:38:23	4
14	830329091	830329091	5000.0	4877.75	2016-12-28T06:43:01	4
15	574788567	574788567	2500.0	2500.00	2016-01-02T11:19:46	3
16	574788567	574788567	2500.0	2469.92	2016-01-16T01:01:27	4
17	574788567	574788567	2500.0	2428.67	2016-01-26T14:04:22	
18	574788567	574788567	2500.0	2428.67	2016-01-29T07:17:39	15
19	574788567	574788567	2500.0	2304.46	2016-01-29T07:33:15	19
20	574788567	574788567	2500.0	2108.39	2016-01-29T21:44:33	
21	574788567	574788567	2500.0	2500.00	2016-02-06T08:16:46	10
22	574788567	574788567	2500.0	2391.14	2016-02-12T03:47:24	2
23	574788567	574788567	2500.0	2362.91	2016-02-22T17:32:13	.
24	574788567	574788567	2500.0	2336.74	2016-02-28T15:53:52	2

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transactionAm
25	574788567	574788567	2500.0	2121.58	2016-02-28T16:43:46	✓
26	574788567	574788567	2500.0	2500.00	2016-03-02T21:49:24	✓
27	574788567	574788567	2500.0	2464.86	2016-03-05T22:24:50	2
28	574788567	574788567	2500.0	2440.42	2016-03-09T14:41:15	13
29	574788567	574788567	2500.0	2300.98	2016-03-10T00:59:51	23
30	574788567	574788567	2500.0	2065.24	2016-03-14T06:24:48	15
31	574788567	574788567	2500.0	2500.00	2016-04-01T20:08:33	3
32	574788567	574788567	2500.0	2465.16	2016-04-05T21:44:57	30
33	574788567	574788567	2500.0	2160.23	2016-04-14T05:00:43	16
34	574788567	574788567	2500.0	1998.84	2016-04-26T04:33:33	6
35	574788567	574788567	2500.0	1930.18	2016-04-28T08:08:33	6
36	574788567	574788567	2500.0	2500.00	2016-05-03T21:11:14	4
37	574788567	574788567	2500.0	2456.36	2016-05-12T00:45:51	4
38	574788567	574788567	2500.0	2416.11	2016-05-24T01:35:33	2
39	574788567	574788567	2500.0	2200.98	2016-05-24T01:38:03	2
40	574788567	574788567	2500.0	2416.11	2016-05-24T19:15:52	28
41	574788567	574788567	2500.0	2129.60	2016-05-26T14:32:39	3
42	574788567	574788567	2500.0	1811.65	2016-05-28T04:42:54	
43	574788567	574788567	2500.0	2500.00	2016-06-04T18:45:39	
44	574788567	574788567	2500.0	2495.54	2016-06-11T04:40:06	14
45	574788567	574788567	2500.0	2355.04	2016-06-15T00:52:10	18
46	574788567	574788567	2500.0	2500.00	2016-07-01T14:55:32	
47	574788567	574788567	2500.0	2491.80	2016-07-03T18:33:35	20
48	574788567	574788567	2500.0	2288.12	2016-07-06T13:08:53	
49	574788567	574788567	2500.0	2283.66	2016-07-16T10:28:25	2

```
In [19]: df.describe()
```

```
Out[19]:
```

	creditLimit	availableMoney	transactionAmount	currentBalance
<b>count</b>	786363.000000	786363.000000	786363.000000	786363.000000
<b>mean</b>	10759.464459	6250.725369	136.985791	4508.739089
<b>std</b>	11636.174890	8880.783989	147.725569	6457.442068
<b>min</b>	250.000000	-1005.630000	0.000000	0.000000
<b>25%</b>	5000.000000	1077.420000	33.650000	689.910000
<b>50%</b>	7500.000000	3184.860000	87.900000	2451.760000
<b>75%</b>	15000.000000	7500.000000	191.480000	5291.095000
<b>max</b>	50000.000000	50000.000000	2011.540000	47498.810000

## Question 1: Load

Programmatically download and load into your favorite analytical tool the transactions data. This data, which is in line-delimited JSON format, can be found [here](#)


**Please describe the structure of the data. Number of records and fields in each record?**

There are 786363 transactions in the JSON data. It is list of dictionaries. There are a total of 786363 transaction records. For each record, there are 29 fields.


**Please provide some additional basic summary statistics for each field. Be sure to include a count of null, minimum, maximum, and unique values where appropriate.**

1. There are 4562 missing data in acqCountry, 724 in merchantCountryCode, 4054 in posEntryMode, 409 in posConditionCode, 698 in transactionType. There are no data for these 6 categories: echoBuffer, merchantCity, merchantState, merchantZip, posOnPremises, recurringAuthInd.
2. The range for creditLimit is from 250 to 50000. The range for availableMoney is from -1005.63 to 50000. The range for transactionAmount is from 0 to 2011.54. The range for currentBalance is from 0 to 47498.81.
3. There are a total of 5000 unique customer ID in this dataset. For these customers, there are a total of 5 countries recorded.

```
In [22]: hist = df.transactionAmount.hist(bins=10)
```




```
In [23]: ax = df.transactionAmount.plot.hist(bins=15, alpha=0.5)
```



## Question 2: Plot

Plot a histogram of the processed amounts of each transaction, the transactionAmount column.



Report any structure you find and any hypotheses you have about that structure.

Exponential decay.

Number of transations decrease exponentially as the transaction amount increases

```
In [24]: # identification of reverse transaction :
id=df[df.transactionType=='REVERSAL'].index
reverse_df=df2.iloc[id,:]
reverse_df
```

```
Out[24]:
```

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transactio
<b>39</b>	574788567	574788567	2500.0	2200.98	2016-05-24T01:38:03	
<b>73</b>	574788567	574788567	2500.0	2432.72	2016-10-07T10:23:57	
<b>101</b>	924729945	924729945	50000.0	49831.43	2016-10-19T14:01:45	
<b>133</b>	984504651	984504651	50000.0	46367.41	2016-01-16T09:53:15	
<b>156</b>	984504651	984504651	50000.0	41909.30	2016-01-25T20:39:15	
...	...	...	...	...	...	
<b>786106</b>	899818521	899818521	2500.0	968.33	2016-09-29T02:04:32	
<b>786120</b>	638498773	638498773	10000.0	9798.21	2016-01-01T19:48:03	
<b>786219</b>	638498773	638498773	10000.0	5331.33	2016-11-03T04:23:26	
<b>786225</b>	638498773	638498773	10000.0	4393.10	2016-11-06T22:54:25	
<b>786301</b>	732852505	732852505	50000.0	49860.23	2016-06-22T19:07:55	

20303 rows × 19 columns

```
In [25]: reverse_df.shape
```

```
Out[25]: (20303, 19)
```

```
In [26]: reverse_df.transactionAmount.sum()
```

```
Out[26]: 2821792.5
```

```
In [27]: from datetime import datetime
```

```
In [34]: df2['transactionDateTime']=pd.to_datetime(df2['transactionDateTime'],format='%Y-
```

```
In [35]: df2['transactionDateTime']
```

```
Out[35]: 0      2016-08-13 14:27:32
1      2016-10-11 05:05:54
2      2016-11-08 09:18:39
```



```
3      2016-12-10 02:14:50
4      2016-03-24 21:04:46
```

```
...
```

```
786358 2016-12-22 18:44:12
786359 2016-12-25 16:20:34
786360 2016-12-27 15:46:24
786361 2016-12-29 00:30:55
786362 2016-12-30 20:10:29
```

```
Name: transactionDateTime, Length: 786363, dtype: datetime64[ns]
```

```
In [36]: customers = df.customerId.unique()
```

```
In [37]: # initialize dataframe
Multi_df=df2.iloc[1:2,:] #all multi-swipe transaction
Multi_df2=df2.iloc[1:2,:] #multi-swipe transactions exclude the 1st normal payment
type(Multi_df)
Multi_df
```

```
Out[37]:
```

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transactionAmount
1	737265056	737265056	5000.0	5000.0	2016-10-11 05:05:54	74

```
In [38]: # multiple swipe: multiple transaction of the same amount, same customer, same merchant
for customer in customers:
    df_cus=df2[df2.customerId==customer]
    for i in range(df_cus.shape[0]-1):
        dff=df2['transactionDateTime'][i+1]-df2['transactionDateTime'][i]
        diff=dff.total_seconds()
        if df_cus.iloc[i,5]==df_cus.iloc[i+1,5] and df_cus.iloc[i,6]==df_cus.iloc[i+1,6]:
            Multi_df=Multi_df.append(df_cus.iloc[i,:])
            Multi_df=Multi_df.append(df_cus.iloc[i+1,:])
            Multi_df2=Multi_df2.append(df_cus.iloc[i+1,:]) #multi-swipe transaction
```

```
In [41]: Multi_df_f=Multi_df.copy()
Multi_df_f=Multi_df_f.drop(1)
```

```
In [42]: #all multi-swipe transaction
Multi_df_f.head(20)
```

```
Out[42]:
```

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transactionAmount
3068	101380713	101380713	10000.0	2407.85	2016-11-17 14:32:54	
3069	101380713	101380713	10000.0	2000.18	2016-11-17 14:35:32	
3891	419989841	419989841	5000.0	831.15	2016-03-04 21:26:00	
3892	419989841	419989841	5000.0	517.56	2016-03-04 21:26:53	
4400	281639186	281639186	2500.0	2477.27	2016-06-11 21:55:24	
4401	281639186	281639186	2500.0	2495.20	2016-07-13 01:22:56	
4649	245118458	245118458	15000.0	15000.00	2016-07-04 21:34:49	

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transaction,
<b>4650</b>	245118458	245118458	15000.0	14921.21	2016-07-04 21:37:14	
<b>4902</b>	288118894	288118894	7500.0	7476.05	2016-04-08 02:48:38	
<b>4903</b>	288118894	288118894	7500.0	7493.66	2016-05-08 08:13:42	
<b>5040</b>	988172671	988172671	250.0	213.05	2016-12-02 23:21:34	
<b>5041</b>	988172671	988172671	250.0	189.28	2016-12-16 05:47:41	
<b>5631</b>	935981871	935981871	5000.0	4897.37	2016-04-06 00:38:27	
<b>5632</b>	935981871	935981871	5000.0	4914.29	2016-05-06 19:09:23	
<b>6689</b>	996362843	996362843	1000.0	977.03	2016-03-04 23:21:04	
<b>6690</b>	996362843	996362843	1000.0	878.66	2016-03-04 23:22:29	
<b>8097</b>	687365478	687365478	10000.0	10000.00	2016-10-02 00:15:22	
<b>8098</b>	687365478	687365478	10000.0	10000.00	2016-10-16 11:09:02	
<b>9359</b>	717714059	717714059	15000.0	15000.00	2016-10-08 16:29:31	
<b>9360</b>	717714059	717714059	15000.0	14964.06	2016-10-22 00:01:37	

In [43]: `Multi_df_f.shape`

Out[43]: (1332, 19)

In [44]: `Multi_df_f2=Multi_df2.copy()  
Multi_df_f2=Multi_df_f2.drop(1)`

In [45]: `#multi-swipe transactions exclude the 1st normal payment  
Multi_df_f2.shape`

Out[45]: (666, 19)

In [46]: `Multi_df_f2.head(50)`

Out[46]:

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transaction,
<b>3069</b>	101380713	101380713	10000.0	2000.18	2016-11-17 14:35:32	
<b>3892</b>	419989841	419989841	5000.0	517.56	2016-03-04 21:26:53	
<b>4401</b>	281639186	281639186	2500.0	2495.20	2016-07-13 01:22:56	
<b>4650</b>	245118458	245118458	15000.0	14921.21	2016-07-04 21:37:14	
<b>4903</b>	288118894	288118894	7500.0	7493.66	2016-05-08 08:13:42	
<b>5041</b>	988172671	988172671	250.0	189.28	2016-12-16 05:47:41	
<b>5632</b>	935981871	935981871	5000.0	4914.29	2016-05-06 19:09:23	
<b>6690</b>	996362843	996362843	1000.0	878.66	2016-03-04 23:22:29	

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transaction
8098	687365478	687365478	10000.0	10000.00	2016-10-16 11:09:02	
9360	717714059	717714059	15000.0	14964.06	2016-10-22 00:01:37	
10525	574212417	574212417	500.0	500.00	2016-11-25 00:12:28	
10922	386732203	386732203	5000.0	1277.50	2016-11-02 23:42:12	
11098	222122504	222122504	15000.0	15000.00	2016-05-19 00:51:02	
13206	393869032	393869032	15000.0	7356.99	2016-11-09 06:48:19	
13375	378369270	378369270	10000.0	10000.00	2016-11-14 15:17:09	
15237	111113489	111113489	15000.0	10954.51	2016-09-21 14:42:10	
15350	364197707	364197707	7500.0	6420.60	2016-11-14 10:53:10	
16669	465894701	465894701	250.0	162.02	2016-11-03 21:27:35	
19082	897665697	897665697	2500.0	2295.20	2016-03-15 21:54:52	
19093	897665697	897665697	2500.0	1790.53	2016-07-04 22:28:07	
19323	543149097	543149097	5000.0	5000.00	2016-09-15 19:35:52	
19453	360564098	360564098	20000.0	19816.50	2016-06-27 23:10:51	
19621	922209733	922209733	2500.0	671.02	2016-10-09 12:19:05	
21110	294270581	294270581	5000.0	5000.00	2016-07-02 05:32:22	
21523	835482161	835482161	7500.0	6359.35	2016-01-19 01:06:04	
22579	368960526	368960526	5000.0	4602.97	2016-12-06 16:25:37	
23650	509526607	509526607	500.0	500.00	2016-09-19 11:20:44	
23896	237944130	237944130	15000.0	14716.52	2016-07-25 11:58:09	
24001	371832344	371832344	50000.0	49985.65	2016-07-04 05:53:01	
24421	722665134	722665134	500.0	330.38	2016-10-26 13:09:48	
24509	605042720	605042720	5000.0	5000.00	2016-07-28 03:05:13	
25004	937999128	937999128	7500.0	6556.23	2016-12-11 06:14:04	
25241	576620175	576620175	5000.0	5000.00	2016-08-27 17:25:36	
25550	888097952	888097952	20000.0	19948.02	2016-06-24 16:01:18	
25561	888097952	888097952	20000.0	19840.42	2016-10-08 21:38:04	
25647	869071202	869071202	500.0	158.92	2016-08-06 23:58:46	
27686	783013155	783013155	500.0	159.09	2016-03-17 03:35:50	
29537	630456222	630456222	250.0	213.10	2016-08-11 05:38:20	
29546	192118573	192118573	7500.0	7253.91	2016-01-02 20:15:15	
29881	192118573	192118573	7500.0	2460.40	2016-05-15 09:34:59	
31221	162576842	162576842	2500.0	2500.00	2016-03-26 05:32:46	
31555	262032771	262032771	7500.0	7138.65	2016-06-05 16:39:24	
31694	739885668	739885668	500.0	292.17	2016-05-27 09:49:14	

	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transaction
<b>32447</b>	985848672	985848672	5000.0	1666.85	2016-03-02 14:27:48	
<b>33038</b>	593072186	593072186	5000.0	5000.00	2016-06-03 18:24:20	
<b>33327</b>	626983390	626983390	10000.0	9924.03	2016-10-10 01:06:33	
<b>33338</b>	626983390	626983390	10000.0	9772.09	2016-12-25 18:18:25	
<b>35949</b>	256583918	256583918	10000.0	4506.16	2016-04-19 04:01:42	
<b>36439</b>	967020232	967020232	20000.0	20000.00	2016-11-15 10:47:31	
<b>45874</b>	410523603	410523603	5000.0	2800.78	2016-11-29 08:19:12	

```
In [47]: Multi_df_f2.transactionAmount.sum()
```

```
Out[47]: 40266.25
```

```
In [48]: Multi_df_f2.merchantCategoryCode.unique()
```

```
Out[48]: array(['rideshare', 'online_retail', 'mobileapps', 'food', 'fuel',
                'fastfood', 'furniture', 'entertainment', 'food_delivery',
                'hotels', 'health', 'online_subscriptions', 'personal care',
                'airline', 'gym', 'subscriptions', 'cable/phone', 'online_gifts',
                'auto'], dtype=object)
```

```
In [49]: # Check for merchantCategoryCode in different transactions
all_pie=df2.groupby(by=["merchantCategoryCode"])[ "merchantCategoryCode"].count()
reverse_pie=reverse_df.groupby(by=["merchantCategoryCode"])[ "merchantCategoryCode"]
multi_pie=Multi_df_f2.groupby(by=["merchantCategoryCode"])[ "merchantCategoryCode"]


print(len(all_pie))
print(len(reverse_pie))
print(len(multi_pie))
```

```
19
13
19
```

```
In [50]: import matplotlib.pyplot as plt
# Pie chart, where the slices will be ordered and plotted counter-clockwise:
# Check for merchantCategoryCode in all transactions
labels = all_pie.index
sizes = all_pie.values
explode = (0.1, 0.1, 0.1, 0.1, 0.1,0.1,0.1,0.1,0.1,0.1,0.1,0.1,0.1,0.1, 0

fig1, ax1 = plt.subplots()
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%',
        shadow=True, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.

plt.show()
```




In [51]:

```
# Pie chart, where the slices will be ordered and plotted counter-clockwise:
# Check for merchantCategoryCode in reversed transactions
labels = reverse_pie.index
sizes = reverse_pie.values
explode = (0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1) # onl


fig1, ax1 = plt.subplots()
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%',
        shadow=True, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.

plt.show()
```



In [52]:


```
#check distribution of transaction amount in reveresd transaction
ax = reverse_df.transactionAmount.plot.hist(bins=15, alpha=0.5)
```



```
In [53]: # Pie chart, where the slices will be ordered and plotted counter-clockwise:
# Check for merchantCategoryCode in multi-swiped transactions
labels = multi_pie.index
sizes = multi_pie.values
explode = (0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1, 0.1) # only "explode" the 2nd slice (i.e. 'Hogs')


fig1, ax1 = plt.subplots()
ax1.pie(sizes, explode=explode, labels=labels, autopct='%1.1f%%',
        shadow=True, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.

plt.show()
```



```
In [54]: #check distribution of transaction amount in multiswiped transaction

ax = Multi_df_f2.transactionAmount.plot.hist(bins=15, alpha=0.5)
```



```
In [55]: sorted(multi_pie.index)
```

```
Out[55]: ['airline',  
          'auto',  
          'cable/phone',  
          'entertainment',  
          'fastfood',  
          'food',  
          'food_delivery',  
          'fuel',  
          'furniture',  
          'gym',  
          'health',  
          'hotels',  
          'mobileapps',  
          'online_gifts',  
          'online_retail',  
          'online_subscriptions',  
          'personal care',  
          'rideshare',  
          'subscriptions']
```

```
In [56]: sorted(reverse_pie.index)
```

```
Out[56]: ['airline',  
          'auto',  
          'entertainment',  
          'fastfood',  
          'food',  
          'furniture',  
          'health',  
          'hotels',  
          'online_gifts',  
          'online_retail',  
          'personal care',  
          'rideshare',  
          'subscriptions']
```

```
In [57]: Rev_id=reverse_df.index  
Rev_id  
Multi_id=Multi_df_f2.index  
Multi_id  
bad_id=Rev_id.append(Multi_id)
```

```
In [58]: df_clean=df2.drop(bad_id)
df_clean.shape
```

```
Out[58]: (765394, 19)
```

## Question 3: Data Wrangling - Duplicate Transactions

You will notice a number of what look like duplicated transactions in the data set. One type of duplicated transaction is a reversed transaction, where a purchase is followed by a reversal. Another example is a multi-swipe, where a vendor accidentally charges a customer's card multiple times within a short time span.

### Can you programmatically identify reversed and multi-swipe transactions?

Reversed transactions are recorded in reverse\_df. They are identified as they have transactionType 'REVERSAL'.

Multi-swipe transactions are recorded in Multi\_df\_f2. They are identified as multiple trasaction of the same amount, same customer, same merchant, within 2 minutes.

**What total number of transactions and total dollar amount do you estimate for the reversed transactions? For the multi-swipe transactions? (please consider the first transaction to be "normal" and exclude it from the number of transaction and dollar amount counts)**

There are a total of 20303 reversed transactions. The total amount for the reversed transactions is 2821792.5.

There are a total of 666 multi-swipe transactions. The total amount for the multi-swipe transactions is 40266.25.

### Did you find anything interesting about either kind of transaction?

There are a total of 19 type of merchant\_Category. Online retail are the merchant\_Category that most subject to mistakes.

For reversed transactions, there are only 13 types involved. The most common type is online retail (28%) and fastfood (16%). There are no following type of reversed transcation: 'cable/phone', 'food\_delivery', 'fuel', 'gym', 'mobileapps', 'online\_subscriptions'. Also, there are higher propotion of transaction in the range of 0-200 compard to that of multi-swipe.

For multi-swipe transactions, all 19 types involved. The most common type is online mobileApps (24%) and online subscription (17%). The two most common type are both online payments.

```
In [59]: #get rid of unwanted/empty columns
df_clean.head()
```

```
Out[59]:  accountNumber  customerId  creditLimit  availableMoney  transactionDateTime  transactionAmo
```



	accountNumber	customerId	creditLimit	availableMoney	transactionDateTime	transactionAmo
0	737265056	737265056	5000.0	5000.0	2016-08-13 14:27:32	98
1	737265056	737265056	5000.0	5000.0	2016-10-11 05:05:54	74
2	737265056	737265056	5000.0	5000.0	2016-11-08 09:18:39	7
3	737265056	737265056	5000.0	5000.0	2016-12-10 02:14:50	7
4	830329091	830329091	5000.0	5000.0	2016-03-24 21:04:46	7

In [60]:

```
df_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 765394 entries, 0 to 786362
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   accountNumber                        765394 non-null object
1   customerId                          765394 non-null object
2   creditLimit                         765394 non-null float64
3   availableMoney                     765394 non-null float64
4   transactionDateTime                 765394 non-null datetime64[ns]
5   transactionAmount                  765394 non-null float64
6   acqCountry                         760958 non-null object
7   merchantCountryCode                764690 non-null object
8   merchantCategoryCode               765394 non-null object
9   currentExpDate                     765394 non-null object
10  accountOpenDate                    765394 non-null object
11  dateOfLastAddressChange             765394 non-null object
12  cardCVV                            765394 non-null object
13  enteredCVV                         765394 non-null object
14  transactionType                     764696 non-null object
15  currentBalance                      765394 non-null float64
16  cardPresent                         765394 non-null bool
17  expirationDateKeyInMatch            765394 non-null bool
18  isFraud                            765394 non-null bool
dtypes: bool(3), datetime64[ns](1), float64(4), object(11)
memory usage: 101.5+ MB
```

In [101]...

```
#get rid of unwanted/empty columns based on later analysis

df_model=df_clean[ ['#creditLimit',
                    'availableMoney',
                    #'transactionDateTime',
                    'merchantCountryCode',
                    'transactionAmount',
                    'merchantCategoryCode','transactionType'
                    , 'currentBalance','cardPresent','expirationDateKeyInMatch','is
```

In [102]...

```
#df_model['currentExpDate']=pd.to_datetime(df_model['currentExpDate'],format='%m
#df_model.dateOfLastAddressChange=pd.to_datetime(df_model['dateOfLastAddressChan
#df_model.accountOpenDate=pd.to_datetime(df_model['accountOpenDate'],format='%Y-
```

In [62]:

```
## I was going to check if months and hour affect the fraud, but turns out these
## but increased computation significantly.
```

```
#df_model['month']=0
#df_model['hour']=0
#for i in range(df_model.shape[0]):
#    dt=df_model.iloc[i,2]
#    df_model['month']= dt.month
#    df_model['hour']= dt.hour
```

<ipython-input-62-679d9187439d>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_model['month']=0
```

<ipython-input-62-679d9187439d>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_model['hour']=0
```

<ipython-input-62-679d9187439d>:5: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_model['month']= dt.month
```

<ipython-input-62-679d9187439d>:6: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df_model['hour']= dt.hour
```

In [103...

```
#df_model.month=df_model.month.astype(str)
#df_model.hour=df_model.hour.astype(str)
df_model.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 765394 entries, 0 to 786362
```

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	availableMoney	765394 non-null	float64
1	merchantCountryCode	764690 non-null	object
2	transactionAmount	765394 non-null	float64
3	merchantCategoryCode	765394 non-null	object
4	transactionType	764696 non-null	object
5	currentBalance	765394 non-null	float64
6	cardPresent	765394 non-null	bool
7	expirationDateKeyInMatch	765394 non-null	bool
8	isFraud	765394 non-null	bool

```
dtypes: bool(3), float64(3), object(3)
```

```
memory usage: 59.2+ MB
```

In [104...

```
# creating a copy of the original data frame
df3 = df_model.copy()

# calling the get_dummies method returns the dummies for all categorical columns
df3 = pd.get_dummies(df_model,
```

```

columns = ['merchantCategoryCode', 'merchantCountryCode', 't
df3=df3.drop(['cardPresent_False', 'expirationDateKeyInMatch_False', 'isFraud_Fals
display(df3)

```

	availableMoney	transactionAmount	currentBalance	merchantCategoryCode_airline	merc
0	5000.00	98.55	0.00		0
1	5000.00	74.51	0.00		0
2	5000.00	7.47	0.00		0
3	5000.00	7.47	0.00		0
4	5000.00	71.18	0.00		0
...	...	...	...		...
786358	48904.96	119.92	1095.04		0
786359	48785.04	18.89	1214.96		0
786360	48766.15	49.43	1233.85		0
786361	48716.72	49.89	1283.28		0
786362	48666.83	72.18	1333.17		0

765394 rows x 31 columns

In [105...

```

from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
from matplotlib import pyplot


```

In [106...

```

#check the fraud outcome distribution
import seaborn as sns
ax=sns.countplot(x=df3['isFraud_True'])
plt.show()

```



Logistic regression to check importance of parameters

In [107...

```

# define dataset

```

```
#y = df['incident_diabetes'][0:500]
y = df3['isFraud_True']
#X=ndf.values[0:500,]
#X=ndf.iloc[0:500,]
X=df3.drop(['isFraud_True'],axis=1)
print(len(y))
X.shape
```

765394

Out[107... (765394, 30)

```
In [108... # define the model
model = LogisticRegression()
# fit the model
res=model.fit(X, y)
```

```
In [109... # get importance
importance = model.coef_[0]
```

```
In [110... FeatureImportance=pd.DataFrame(zip(X.columns,np.transpose(model.coef_.tolist())[0
FeatureImportance['abs_importance']=abs(FeatureImportance.coef)
FeatureImportance.sort_values(by=['abs_importance'],ascending=False)
```

```
Out[110...
           features      coef  abs_importance
25  merchantCountryCode_US -1.442520      1.442520
27  transactionType_PURCHASE -1.353383      1.353383
28  cardPresent_True -0.718365      0.718365
17  merchantCategoryCode_online_retail -0.264183      0.264183
7   merchantCategoryCode_fastfood -0.222783      0.222783
8   merchantCategoryCode_food -0.167477      0.167477
6   merchantCategoryCode_entertainment -0.161233      0.161233
26  transactionType_ADDRESS_VERIFICATION -0.098699      0.098699
10  merchantCategoryCode_fuel -0.097059      0.097059
16  merchantCategoryCode_online_gifts -0.091584      0.091584
15  merchantCategoryCode_mobileapps -0.077095      0.077095
20  merchantCategoryCode_rideshare -0.065211      0.065211
18  merchantCategoryCode_online_subscriptions -0.056763      0.056763
4   merchantCategoryCode_auto -0.048270      0.048270
14  merchantCategoryCode_hotels -0.045530      0.045530
13  merchantCategoryCode_health -0.039057      0.039057
21  merchantCategoryCode_subscriptions -0.033375      0.033375
19  merchantCategoryCode_personal care -0.032459      0.032459
9   merchantCategoryCode_food_delivery -0.021027      0.021027
```

	features	coef	abs_importance
11	merchantCategoryCode_furniture	-0.011372	0.011372
12	merchantCategoryCode_gym	-0.009674	0.009674
3	merchantCategoryCode_airline	-0.005580	0.005580
23	merchantCountryCode_MEX	-0.005118	0.005118
22	merchantCountryCode_CAN	-0.003955	0.003955
5	merchantCategoryCode_cable/phone	-0.003370	0.003370
1	transactionAmount	0.002780	0.002780
24	merchantCountryCode_PR	-0.002779	0.002779
29	expirationDateKeyInMatch_True	-0.002181	0.002181
2	currentBalance	0.000004	0.000004
0	availableMoney	-0.000004	0.000004

## ML prediction

```
In [111...
import xgboost as xgb
from sklearn.ensemble import AdaBoostClassifier
from sklearn.datasets import make_classification
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve, auc, recall_score, precision_score, roc_auc_
```

```
In [112...
#use multiple ml algorithms for model fitting
from sklearn.svm import SVC
from sklearn.metrics import plot_roc_curve
from sklearn.ensemble import RandomForestClassifier

X, y = make_classification(random_state=0)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_

svc = SVC(random_state=42)
svc.fit(X_train, y_train)
#rfc = RandomForestClassifier(random_state=42)
rfc = RandomForestClassifier(max_depth=6, random_state=42)
rfc.fit(X_train, y_train)
logreg=LogisticRegression(random_state=42)
logreg.fit(X_train, y_train)
xg_reg = xgb.XGBClassifier(random_state=42)
xg_reg.fit(X_train, y_train)
adareg=AdaBoostClassifier(n_estimators=100,
                           learning_rate=0.5, random_state=42)
adareg.fit(X_train, y_train)
#xg_reg = xgb.XGBClassifier(random_state=42)
xg_reg = xgb.XGBClassifier(objective='binary:logistic', colsample_bytree = 0.1,
                           max_depth = 25, alpha = 10, n_estimators = 300, booster='gbtree')
xg_reg.fit(X_train, y_train)

svc_disp = plot_roc_curve(svc, X_test, y_test)
rfc_disp = plot_roc_curve(rfc, X_test, y_test)
```

```

logreg_disp = plot_roc_curve(logreg, X_test, y_test)
adareg_disp = plot_roc_curve(adareg, X_test, y_test)
xgbreg_disp = plot_roc_curve(xg_reg, X_test, y_test)
#xgb_disp.figure_.suptitle("ROC curve comparison")

plt.show()

```

/Users/ziye/opt/anaconda3/lib/python3.8/site-packages/xgboost/sklearn.py:888: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use\_label\_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num\_class - 1].


```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```




/Users/ziye/opt/anaconda3/lib/python3.8/site-packages/xgboost/sklearn.py:888: UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use\_label\_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num\_class - 1].


```
warnings.warn(label_encoder_deprecation_msg, UserWarning)
```

[04:28:00] WARNING: /opt/concourse/worker/volumes/live/7a2b9f41-3287-451b-6691-43e9a6c0910f/volume/xgboost-split\_1619728204606/work/src/learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval\_metric if you'd like to restore the old behavior.

[04:28:00] WARNING: /opt/concourse/worker/volumes/live/7a2b9f41-3287-451b-6691-43e9a6c0910f/volume/xgboost-split\_1619728204606/work/src/learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval\_metric if you'd like to restore the old behavior.







In [120...

```

svcpreds = svc.predict(X_test)
accuracy=accuracy_score(y_test,svcpreds)
print('The accuracy for SVC is {}'.format(accuracy))

rfcpreds = rfc.predict(X_test)
accuracy=accuracy_score(y_test,rfcpreds)
print('The accuracy for random forest is {}'.format(accuracy))

logregpreds = logreg.predict(X_test)
accuracy=accuracy_score(y_test,logregpreds)
print('The accuracy for logistic regression is {}'.format(accuracy))

adapreds = adareg.predict(X_test)
accuracy=accuracy_score(y_test,adapreds)
print('The accuracy for ada boost is {}'.format(accuracy))

xgpreds = xg_reg.predict(X_test)
accuracy=accuracy_score(y_test,xgpreds)
print('The accuracy for xgboost is {}'.format(accuracy))

```

```

The accuracy for SVC is 0.9
The accuracy for random forest is 0.95
The accuracy for logistic regression is 0.8
The accuracy for ada boost is 0.9
The accuracy for xgboost is 0.95

```

## Question 4: Model

Fraud is a problem for any bank. Fraud can take many forms, whether it is someone stealing a single credit card, to large batches of stolen credit card numbers being used on the web, or even a mass compromise of credit card numbers stolen from a merchant via tools like credit card skimming devices.


Each of the transactions in the dataset has a field called isFraud. Please build a predictive model to determine whether a given transaction will be fraudulent or not. Use as much of the data as you like (or all of it).

**Provide an estimate of performance using an appropriate sample, and show your work.**



Because a very small proportion of data is positive, I used AUCROC score to evaluate the performance. The data use split at 0.8/0.2 for training/test. Comparing the performance of SVC, Random forest, logistic regression, adaboost and xgboost, it is found that xgboost gives the best AUC of 0.99. The accuracy of xgboost is also the highest among algorithms.

Please explain your methodology (modeling algorithm/method used and why, what features/data you found useful, what questions you have, and what you would do next with more time)



1. Since this is a binary prediction for fraud, I used binary Classification models including SVC, Random forest logistic regression, adaboost and xgboost, it is found that xgboost gives both the best AUC score of 0.99 and the best accuracy of 0.95.
2. Using a parameter analysis on feature importance, the creditLimit data was excluded due to low feature importance. Some of the top important features are: merchantCountryCode\_US, cardPresent\_True, merchantCategoryCode\_online\_retail. The following variable is included: 'availableMoney', 'transactionDateTime' (transformed to months and hours), 'merchantCountryCode', 'transactionAmount', 'merchantCategoryCode', 'transactionType', 'currentBalance', 'cardPresent', 'expirationDateKeyInMatch', 'isFraud'
3. Because there are much more non-fraud cases compared to fraud cases, the more abundant categories seem to have a greater negative impact. A more balanced dataset, i.e

a greater proportion of fraud cases included, may change the feature importance and



increase prediction rate.

4. If given more time, I wish to perform clustering analysis to cluster customers into groups by features, and see if each groups have different fraud rate and prediction accuracy.

In [ ]: